

A2 · NOTICE TO ALL TEAMS

What changed in the A2 documents

Posted Tuesday 1 September 2026 · covers every change since Monday afternoon

Affects	All four A2 documents	Action	Re-download all four
New document	TEAM_DECLARATION — due Fri 4 Sep	Deadline	Unchanged — Sun 13 Sep

I have made a round of corrections and clarifications to the A2 documents. Most are things you asked about, a few are errors I found and fixed, and one document is entirely new.

Please re-download all four files. If you started work over the weekend, nothing you have done is wasted. No deliverable was removed, no deliverable was added, and no deadline moved except the peer rating, which moved one day *earlier*.

If you only read one thing

Four points carry actual work. Everything else is a correction or a clarification.

1. **The starter material arrives tomorrow.** Reference data and the code scaffold together, by midday **Wednesday 2 September** — not Monday, as three places in the brief used to say.
2. **Every team member now runs one live model.** Three models is the minimum that passes. We expect one model per member, minus the one person who runs the v1 prompt comparison — so five models plus a v1 pass in a team of six.
3. **Only one part of A2 spends money.** The guardrail checklist, the reproducible end-to-end run and the two failure reproductions all run on the scripted backend, free. This is newly spelled out, and it should save you tokens.
4. **There is a new team declaration form, due Friday 4 September.** It did not exist before. Section 8 always asked for the check-in; now there is a form to fill in.

1 · The A2 brief — PE6201_A2_Applied_AI_System.pdf

Dates

- **Starter material now lands midday Wednesday 2 September**, not Monday. Three places said Monday; all now say Wednesday. Nothing in D0, D2(a), D2(b) or D3(a) needs the scaffold, so you were never actually blocked.
- **Peer rating moved to Wednesday 16 September**, 23:59 SGT. It previously read “Wednesday 17 September” — but 17 September 2026 is a Thursday, so one of those two had to be wrong. I kept the weekday and moved the date.
- The brief also said “Monday 1 September” in two places. **1 September 2026 was a Tuesday**. Corrected. There is now an automatic check that compares every printed weekday against the real 2026 calendar before a document can go out.

D2(c) — the worked example on tool calling

- **The turn counts in Problem A's example were wrong, and a student spotted it.** The code block lists eight tool calls, but the text said six turns becoming three. It is **eight sequential turns becoming four**, and the saving is **20,800 → 9,600 tokens, 54%** (it used to print 13,200 → 6,000). The eight calls are now named one by one so you can count them, and the example points at the matching record in Appendix A so the two visibly agree.
- **Problem B now has its own parallel-calling worked example.** Referral REF-5602 — six calls, four turns, **13,200 → 8,400, a 36% saving**. Deliberately a smaller saving than A's, because B's work is a chain (band, then slots, then book) and a chain cannot be shortened by running things at the same time. If you were leaning toward Problem A only because it had the worked example, that reason is gone.
- **New box: these are examples, not the prescribed answer.** If your team parallelises less, groups calls differently, or keeps a call on its own so you can read its result first, you have not made a mistake. **D2(c) marks your reasoning, not your number.** What is required is a stated dependency rule, the measurement taken both ways, and evidence that correctness did not change.

D3 — guardrails

- **New: run the ten guardrail cases on the scripted backend.** A step cap, a budget ceiling, action de-duplication and the autonomy gate are all *your code* — a model cannot influence whether they fire. So a deterministic backend is the correct instrument, not merely the cheap one. **These ten cases cost nothing and need no API key.**
- One caveat is now spelled out. A scripted run proves *your guardrail fires* when the agent attempts a bad action. It cannot tell you whether a live model gets talked into attempting it. That second question belongs to your D5 battery, not to a guardrail case.

D4 — the evaluation set

This section was rewritten. It is the largest change in the brief.

- **It now reads the same whichever problem you chose.** Everything is phrased in terms of the three outcomes — act, ask, escalate — rather than naming claims or referrals.
- **“L1 / L2” is gone.** Several of you found it confusing in A1. There are now two kinds of check, described plainly. A **code check**: your harness compares the answer against your answer key — no model, no person, no opinion. A **judgement check**: a person, *or a second model*, reads the record and decides. Same kind of check; the only difference is who grades. You need both, and there is now one example each for Problem A and Problem B.
- **The run arithmetic is now a table, so nobody has to guess.** Minimum that passes: 30 cases, 6 negative. What we expect: 40 cases, 8 negative. Ordinary cases get **one trial**; **negative cases get three** — two extra each, because negatives are the ones that flip between runs and a single trial cannot tell a real refusal from a lucky one. That is **56 runs per model** at the expected shape.
- **The v1-versus-v2 comparison does not double your battery.** v1 runs on **one** model only — the same model you quote its v2 against. To compare prompt versions you hold the model fixed; to compare models you hold the prompt fixed. Change one thing at a time or the difference is not attributable to anything.
- **The word “policy” is no longer used for prompt versions.** In Problem A “policy” already means the member's insurance policy, which was confusing. It now says **prompt version** throughout.

D5 — the model battery

- The heading is now **“one scripted run, one live model per member”**.
- **(a) The scripted run is unchanged in substance.** Your submitted harness must run end to end on a scripted backend, with `BACKEND = "scripted"` as the default, so a marker can clone your repository and reproduce your numbers at zero cost. **Technical Execution is capped if it does not run this way.**
- **(b) New rule on how many live models.** Three is the floor that passes. **$N - 1$ models for a team of N is what we expect**, with the remaining member running the v1 pass. Five models plus a v1 pass in a team of six; six plus a v1 pass in a team of seven. Every member runs one full battery on their own key.
- **This costs nobody anything extra.** The price table in section 7 is *per member, per model*. Fifty-six runs on the cheap tier is about US\$0.27 whether your team fields three models or six.
- **Two conditions, or the comparison is meaningless.** Your models must span **at least two price tiers**, and no two members may pick models from the same family. And every member must run the **identical evaluation set and identical v2 prompt** — the model name is the only thing that may differ between you.
- **Writing it up.** Report section 3 is still 350 words. Tables do not count toward the word cap, so put every model in the results table and spend the prose on **the spread and the two extremes** — the cheapest model that met your bar, the most expensive one that did not earn its price, and where the negative cases separated them.

How you divide the rest of the work is your team's decision. The agent, the tools, the guardrail layer, the harness, the scripted run and the two failure reproductions are **the build** — two weeks of real engineering and most of the marks. We are not prescribing who does which part of it. Agree it at your first team meeting and record the owners in section 4 of the team declaration.

Only two things are fixed: everyone writes evaluation cases, and everyone runs one live model at the end.

D7 — the two reproduced failures

- **New: both failures and their before/after tables run on the scripted backend.** A failure built as a deletion from your working agent is deterministic by construction — that is the point of building it that way. Script the observation that causes the loop, and the loop reproduces for ever at no cost. **D7 needs no key, and it must reproduce for a marker.**

Section 7 — the budget

- **The cost table was wrong and is now corrected.** It was pricing 64 runs by accidentally giving negative cases four trials instead of three. The correct figures for one member running one full 56-run battery are **US\$0.27 cheap tier, US\$2.76 mid, US\$13.78 frontier** (they used to read 0.31 / 3.15 / 15.74).
- Read that last figure again: **one member on a frontier model spends more than the whole course allowance.** If you want a frontier model in your comparison, run it on the **negative cases only** and say so in the report.
- **“Three rules that keep you inside the budget” is now four.** The new rule names D3(b), D5(a) and D7 as scripted and therefore free. If you find yourself spending live tokens on any of the three, stop.

Appendix A and the glossary

- **Tool names are suggestions.** The names in Appendix A describe work that has to happen; they are not an interface for you to implement. Rename, re-argument, merge, split or add tools as your design requires. The reason is D2(b): you are asked for a six-field descriptor per tool, and you cannot meaningfully write one for a signature somebody else fixed.
- **What is not negotiable is the routing rule and the gated action,** because the answer key is written against them. That sentence has been simplified.
- **Glossary. Routing rule** added — the brief used the term without ever defining it. **Gated action** rewritten in plainer language, with one worked example for Problem A and one for Problem B.
- A join example in section 2 used a tool called `lookup_member` that appears in neither problem's tool set. It is rewritten around `lookup_policy`, which does exist.

2 · The FAQ — PE6201_A2_FAQ.pdf

- **“Applications” removed.** The fixture-data answer referred to “claims / applications / referrals”. *Applications* was the primary record of a problem that was cut from this assignment before release.
- **“30–60 primary records” corrected to 30–50,** which is what the brief says in five other places and what the FAQ's own table said two paragraphs below it.
- **The starter-material question now says Wednesday,** including in the question index at the front.
- **The negative-case table now gives a Problem A example and a Problem B example,** instead of only one problem's.
- **“L1 / L2” replaced with code check and judgement check,** matching the brief.
- **A worked pass-rate example added** — 56 trials, 48 passed — so you can see how the number is reported together with its trial count. **A pass rate quoted without its trial count is not a measurement.**
- **New: how the model battery splits across the team.** One member owns one model and runs it on their own key, and every member owns one. Three is the floor; $N - 1$ for a team of N is what we expect, with the remaining member on the v1 pass. Includes the cost-neutrality point and the two conditions above.
- **The suggested work split is now clearly labelled a suggestion.** The build is yours to divide however suits your team's skills; record the owners in the team declaration. Two rows stay fixed: everyone writes evaluation cases, and everyone runs one live model.

3 · The team self-appraisal — PE6201_A2_Team_Self_Appraisal.docx

- **The A2 timeline is now printed on the sheet,** the same as on the team declaration, so you are not cross-referencing the brief while filling it in.
- **The masthead states the deadline plainly:** DUE SUNDAY 13 SEPTEMBER 2026.
- **“Trials per case” renamed “Live runs per model (D4)”** — the old label invited the wrong number.
- **The peer-rating tick box is replaced with a dated reminder.** Peer rating opens on 13 September, which is the day this sheet is due — so that box could never honestly be ticked “yes”.

- **New: a fixture-integrity declaration.** You confirm you added rows rather than editing or deleting the ones we shipped.
- **New row: your v1 → v2 change,** so the prompt rewrite you measured appears on the sheet rather than only in the repository.
- **New row: the live model YOU ran, and your pass rate on it.** If you ran the v1 pass instead, write “v1 pass” and name the model it sat on.

4 · NEW — the team declaration — TEAM_DECLARATION.docx

This document did not exist before. It is due Friday 4 September 2026, 23:59 SGT.

Section 8 of the brief has always asked for a team check-in by that date, but there was no form — so every team would have invented its own. There is now a two-page template. Use it as it is; do not redesign it.

What it asks for:

1. **Team ID, and every member's name and matriculation number.**
2. **Which problem you chose** — A or B.
3. **Your repository URL.**
4. **Who owns what** — a table of work strands against names. The listed strands are a suggestion you are free to replace with your own. Two rows are not optional: **everyone writes evaluation cases**, and **everyone runs one live model** (name the model beside each person).
5. **A contribution statement.** This section is marked NOT OPTIONAL in red. You state one of two things: that all members are contributing, or that there is a problem you want us to know about now rather than on the 13th.
6. **A signature and date.**

The A2 timeline is printed on the form.

Arriving midday tomorrow — Wednesday 2 September

On NTULearn, alongside these four documents:

- **The reference data for both problems** — claims, members, policies and procedures for A; referrals, patients, specialties and clinic slots for B. Ready to use, plus the scripts that generate them.
- **A data guide** explaining what each file contains, the shape of every field, and how the files connect to one another.
- `check_my_data.py` — a script that checks your own additions. When you add cases of your own it tells you if you have referred to something that does not exist, changed a record we shipped, or duplicated an id.
- **The starter scaffold notebook.**

The rule for extending the data is one line: add new rows with new ids; never edit or delete a row we shipped. Everything else follows from it.

What did NOT change

So you can stop looking for it:

- **The deadline is still Sunday 13 September 2026, 23:59 SGT.** The only date that moved is the peer rating.
- **The report is still capped at 2,000 words.** Tables and figures still do not count toward it.
- **The list of deliverables is unchanged** — D0 through D7, exactly as before. Nothing was added and nothing was removed.
- **What you submit is unchanged** — the repository in two places, the report, the self-appraisal, and the video link.
- **The marking weights are unchanged** — Conceptual Understanding 25%, Technical Execution 30%, Reasoning & Justification 25%, Communication & Clarity 20%.
- **The evaluation set is still 30–50 cases** with at least two negative cases, and still at least ten guardrail cases.

Key dates

Date	What
Wed 2 Sep, midday	Reference data, data guide and starter scaffold on NTULearn
Fri 4 Sep, 23:59	Team declaration due
Sun 13 Sep, 23:59	A2 due — repository, report, self-appraisal, video link
Wed 16 Sep, 23:59	Peer rating due (<i>moved from 17 Sep</i>)
Sun 27 Sep, 23:59	End-of-course project due

Questions to the usual channel. **If something in these documents still contradicts something else, tell me.** Two of the corrections above came from students reading carefully, and both were real errors.